

# **Signals and signal processing (2)**

**Suntae Hwang**  
**Kookmin University**

# Unreliable signals (1/2)

---

- Signals were unreliable in earlier version of UNIX (V7)
  - Signals could get lost
    - ◆ a signal could occur and the process would never know about it
  - The action for the signal was reset to its default each time the signal occurred
  - A process had a little control over a signal
    - ◆ can catch or ignore the signal, but can't *block* it
    - ◆ unable to turn a signal off when it didn't want the signal to occur (all it can do was ignore the signal)
    - ◆ can't "prevent the following signals from occurring, but only *remember* if they do occur".

# Unreliable signals(2/2)

```
int  sig_int(); /*my signal handling function*/
...
signal(SIGINT, sig_int); /*establish handler*/
...
sig_int()
{
    signal(SIGINT, sig_int);
    /*reestablish handler for next occurrence*/
    ...      /*process the signal*/
}
```

```
/* set the flag to remember that an interrupt occurs */
int  sig_int_flag;
main() {
    int sig_int();
    ...
    signal(SIGINT,sig_int);
    ...
    while(sig_int_flag==0)
        pause(); /* go to sleep, wating for signal */
    ...
}
sig_int() {
    signal(SIGINT,sig_int);
    sig_int_flag=1; /*set flag for main loop to
examine*/
}
```

**Q:** What happens if the interrupt signal occurs again before singal() in sig\_int() executes?

**Q:** What happens if the interrupt signal occurs again before pause() and after while statement?

# Reliable Signal Terminology and Semantics

---

- A signal is *generated* for a process (or sent to a process) when the event that causes the signal occurs
  - when signal is generated the kernel usually sets a flag of some form in the process table
- A signal is *delivered* to a process when the action for a signal is taken
- During the time between the generation of signal and its delivery, the signal is said to be *pending*.
- Each process has a *signal mask* that defines the set of signals currently *blocked* from delivery to that process

# Reliable Signal Terminology and Semantics

---

## □ *blocking* the delivery of a signal

- A blocked signal (with default or catch signal action) remains *pending* until
  - ◆ unblocks the signal
  - ◆ changes the action to ignore the signal
- What to do with a blocked signal is determined when the signal is *delivered*, not when it is generated
  - ◆ This allows the process to change the action for the signal before it is delivered
- Most Unix does not queue blocked signals generated more than once; the Unix kernel just delivers the signal once.

# Signal sets

---

- A data type to represent multiple signals (sigset\_t)
- Do not assume global/static variable initialization in C for sigset\_t
- Functions to manipulate signal sets

```
#include <signal.h>

int sigemptyset (sigset_t *set);
int sigfillset (sigset_t *set);
int sigaddset (sigset_t *set, int signo);
int sigdelset (sigset_t *set, int signo);

                                all four return: 0 if OK, -1 on error

int sigismember ( const sigset_t *set, int signo);
                                Returns: 1 if true, 0 if false
```

# Signal Functions

---

- Superset of the functionality of `signal()`
- `sigprocmask(int how, const sigset_t * setp, sigset_t * osetp);`
  - examine or change signal masks
- `sigpending(sigset_t* setp);`
  - Return the set of signals that are blocked and pending
- `sigaction(int signo, const struct sigaction* act, struct sigaction oact);`
  - examine or modify the action associated with a particular signal
- `sigsuspend(const sigset_t* sigmask);`
  - atomic operation (reset the signal mask and pause )
- `sigsetjmp(sigjmp_buf env, int savemask)/siglongjmp(sigjmp_buf env, init val)`
  - nonlocal branching from a signal handler

# Sigprocmask function

```
#include <signal.h>
int sigprocmask(int how, const sigset_t *set, sigset_t *oset);
Returns: 0 if OK, -1 on error
```

- Examine or change the **signal mask** *that is the set of signals currently blocked from delivery to the process*
  - First, if *oset* is nonnull pointer, the current signal mask for the process is returned through *oset*
  - Second, if *set* is a nonnull pointer, then the *how* argument indicates how the current signal mask is modified
  - if *set* is NULL, *how* is not significant
  - *how*
    - ◆ SIG\_BLOCK : redefine the new signal mask (*set* + *oset*)
    - ◆ SIG\_UNBLOCK : unblock the signals in *set*
    - ◆ SIG\_SETMASK : the new signal mask = *set*
- ❖ if there are any pending, unblocked signals after `sigprocmask` at least one of these signal is delivered to the process before `sigprocmask` returns



# Sigprocmask Function Example

```
#include <errno.h>
#include <signal.h>
#include "ourhdr.h"

void
pr_mask(const char *str)
{
    sigset_t sigset;
    int      errno_save;

    errno_save = errno;          /* we can be called by
signal handlers */
    if (sigprocmask(0, NULL, &sigset) < 0)
        err_sys("sigprocmask error");

    printf("%s", str);
    if (sigismember(&sigset, SIGINT)) printf("SIGINT ");
    if (sigismember(&sigset, SIGQUIT)) printf("SIGQUIT ");
    if (sigismember(&sigset, SIGUSR1)) printf("SIGUSR1 ");
    if (sigismember(&sigset, SIGALRM)) printf("SIGALRM ");
    /* remaining signals can go here */
    printf("\n");
    errno = errno_save;
}
```

# Sigpending Function

---

```
#include <signal.h>
int sigpending (sigset_t *set);
```

Returns: 0 if OK, -1 on error

- Returns the set of signals that are blocked from delivery and currently pending for the calling process
  - The set of signals is returned through the *set* argument

# Sigpending Function

```
static void    sig_quit(int);
int main(void) {
    sigset_t newmask, oldmask, pendmask;

    if (signal(SIGQUIT, sig_quit) == SIG_ERR)
        err_sys("can't catch SIGQUIT");
    sigemptyset(&newmask); sigaddset(&newmask, SIGQUIT);

    /* block SIGQUIT and save current signal mask */
    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0)
        err_sys("SIG_BLOCK error");
    sleep(5);                /* SIGQUIT here will remain pending */

    if (sigpending(&pendmask) < 0)
        err_sys("sigpending error");
    if (sigismember(&pendmask, SIGQUIT))
        printf("\nSIGQUIT pending\n");

    /* reset signal mask which unblocks SIGQUIT */
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
        err_sys("SIG_SETMASK error");
    printf("SIGQUIT unblocked\n");
    sleep(5);                /* SIGQUIT here will terminate with core file */
    exit(0);
}
```

# Sigpending Function

```
static void
sig_quit(int signo)
{
    printf("caught SIGQUIT\n");

    if (signal(SIGQUIT, SIG_DFL) == SIG_ERR)
        err_sys("can't reset SIGQUIT");
    return;
}
```

**\$ a.out**

**^\\^\\^\\^\\**

*generate signal several times (before  
5 seconds are up)*

**SIGQUIT pending**

*after return from sleep(5)*

**caught SIGQUIT**

*in signal handler (signal is generated once!)*

**SIGQUIT unblocked**

*after return from sigprocmask*

**^\\Quit(core dump)**

*generate signal again (message by shell)*

# Sigaction Function (1/3)

```
#include <signal.h>
int sigaction (int signo, const struct sigaction *act,
               struct sigaction *oact);
Returns:0 if OK, -1 on error
```

- Examine or modify the action associated with a particular signal -  
supersedes the signal function from earlier UNIX
  - If act pointer is nonnull, we are modifying the action.
  - If oact pointer is nonnull, the system returns the previous action for the signal.
- Specify a set of signals added to the signal mask before the signal handler is called - *includes the signal being delivered*.
  - This way we are able to *block* certain signals whenever a signal handler invoked
  - When the signal-catching function returns, the signal mask of process is *reset* to its previous value
- The action remains installed until we explicitly change it by calling sigaction().

## Sigaction Function (2/3)

```
struct sigaction {  
    void      (*sa_handler) (int);      /* address of signal handler */  
    sigset_t   sa_mask;                 /* additional signals to block */  
    int        sa_flags;                 /* signal options */  
    void (*sa_sigaction) (int, siginfo_t *, void *);  
};
```

- *sa\_handler*: points to the signal handler or SIG\_IGN or SIG\_DFL
- *sa\_mask*: additional signals to block when the signal handler (as opposed to SIG\_IGN or SIG\_DFL) is called
- *sa\_flags* : specifies various options for the handling of the signal
  - SA\_RESTART - system calls interrupted are automatically restarted
  - SA\_NODEFER - when this signal is caught, the signal is not automatically blocked by the system while the signal handler executes (unreliable signal)
  - SA\_RESETHAND - the disposition for the signal is reset to SIG\_DFL on entry to the signal handler (unreliable signal)
  - SA\_SIGINFO – provides additional information to a signal handler. Final *sig\_action* field is used

# Sigation Function (3/3)

```
/* An implementation of reliable signal() using sigaction */
/* for unreliable signal(), use SA_RESETHAND and SA_NODEFER */

#include <signal.h>

Sigfunc *
signal(int signo, Sigfunc *func) {
    struct sigaction    act, oact;

    act.sa_handler = func;
    sigemptyset(&act.sa_mask);
    act.sa_flags = 0;

    if (signo != SIGALRM) {
        act.sa_flags |= SA_RESTART;    /*SVR4, 4.3+BSD*/
    }

    if (sigaction(signo, &act, &oact) < 0)
        return(SIG_ERR);

    return (oact.sa_handler);
}
```

# Sigsetjmp and siglongjmp Functions

```
#include <set jmp.h>
int sigsetjmp (sigjmp_buf env, int savemask) ;
    Returns: 0 if called directly, nonzero if returning from a call to siglongjmp
int siglongjmp (sigjmp_buf env, int val) ;
```

- What happens to the signal mask for the process if we longjmp out of the signal handler?
- sigsetjmp and siglongjmp saves and restores the signal mask - use these functions for nonlocal branching from a signal handler
  - If *savemask* is nonzero then sigsetjmp also saves the current signal mask of the process in *env*
  - When siglongjmp is called, if the *env* argument was saved with nonzero *savemask*, then siglongjmp restores the saved signal mask



# Sigsetjmp and siglongjmp Functions

## <Example of sigsetjmp and siglongjmp>

```
static sigjmp bufjmpbuf;
static volatile sig_atomic_t canjump;
int main(void) {
    ...
    if(sigsetjmp(jmpbuf, 1)) exit(0);
    canjump=1;          /* now sigsetjmp() is OK */
    ...
}
static void sig_usr1(int signo) {
    ...
    if (canjump==0) return; /* unexpected signal, */
    /* not jmpbuf ready, ignored */
    ...
    canjump = 0;
    siglongjmp(jmpbuf, 1);
}
```

## □ Use of *canjump* variable

- protection against the signal handler being called when the jump buffer isn't initialized by sigsetjmp
- sig\_atomic\_t : variable can be written without being interrupted (ex. No page boundary crossing )

# Sigsuspend Function (1/4)

```
#include <signal.h>
int sigsuspend (const sigset_t *sigmask) ;
Returns: -1 with errno set to EINTR
```

- Performs resetting the signal mask and put the process to sleep in a single atomic operation
  - Signal mask of the process is set to the value pointed to by *sigmask*.
  - The process is also suspended until a signal is caught or until a signal occurs that terminates the process
  - If a signal is caught and if the signal handler returns, then *sigsuspend* returns and the signal mask of the process is set to its *old value*

# Sigsuspend Function (2/4)

```
/* protect critical regions of code from interrupt signals : a wrong way */
sigset_t newmask, oldmask;

sigemptyset(&newmask);
sigaddset(&newmask, SIGINT);
    /* block SIGINT and save current signal mask */
if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0)
    err_sys("SIG_BLOCK error");

    /* critical region of code */

    /* reset signal mask which unblocks SIGINT */
if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
    err_sys("SIG_SETMASK error");

pause(); /* wait for signal to occur */

    /* and continue processing ... */
```

□ Problem: any signal between the second sigprocmask() and pause() gets lost.

# Sigsuspend Function (3/4)

```
/* protect critical regions of code from interrupt signals : a right way */
sigset_t  newmask, oldmask, zeromask;
if (signal(SIGINT, sig_int) == SIG_ERR) err_sys("signal(SIGINT) error");

sigemptyset(&zeromask);
sigemptyset(&newmask);
sigaddset(&newmask, SIGINT);
    /* block SIGINT and save current signal mask */
if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0) err_sys("SIG_BLOCK error");

    /* critical region of code */

    /* allow all signals and pause */
if (sigsuspend(&zeromask) != -1) err_sys("sigsuspend error");

    /* reset signal mask which unblocks SIGINT */
if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0) err_sys("SIG_SETMASK error");

    /* and continue processing ... */
```

- Eliminated the previous problems. (unblock and pause)
- Note: need the second `sigprocmask()` to unblock SIGINT because the return of `sigsuspend()` set the signal mask to its value before the call.

# Sigsuspend Function (4/4)

```
volatile sig_atomic_t quitflag;    /* set nonzero by signal handler */
int main(void) {
    sigset_t    newmask, oldmask, zeromask;
    if (signal(SIGINT, sig_int) == SIG_ERR) err_sys("signal(SIGINT) error");
    if (signal(SIGQUIT, sig_int) == SIG_ERR) err_sys("signal(SIGQUIT) error");

    sigemptyset(&zeromask);  sigemptyset(&newmask);
    sigaddset(&newmask, SIGQUIT); /* block SIGQUIT and save current signal mask */

    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0) err_sys("SIG_BLOCK error");

    while (quitflag == 0) sigsuspend(&zeromask);

    /* SIGQUIT has been caught and is now blocked; do whatever */
    quitflag = 0;
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0) err_sys("SIG_SETMASK error");
    exit(0);
}

void sig_int(int signo)  { /* one signal handler for SIGINT and SIGQUIT */
    if (signo == SIGINT) printf("\ninterrupt\n");
    else if (signo == SIGQUIT) quitflag = 1; /* set flag for main loop */
    return;
}
```

□ sigsuspend to wait for a global variable to be set.



# Sleep(1/2)

```
static void sig_alm(void) {
    return; /* nothing to do, just returning wakes up sigsuspend() */
}

unsigned int sleep(unsigned int nsecs) {
    struct sigaction  newact, oldact;
    sigset_t          newmask, oldmask, suspmask;
    unsigned int      unslept;

    newact.sa_handler = sig_alm;
    sigemptyset(&newact.sa_mask);
    newact.sa_flags = 0;
    sigaction(SIGALRM, &newact, &oldact);
                /* set our handler, save previous information */

    sigemptyset(&newmask);
    sigaddset(&newmask, SIGALRM);
                /* block SIGALRM and save current signal mask */
    sigprocmask(SIG_BLOCK, &newmask, &oldmask);
```

## Sleep(2/2)

```
alarm(nsecs);

suspmask = oldmask;
sigdelset(&suspmask, SIGALRM); /* make sure SIGALRM isn't blocked */

sigsuspend(&suspmask);          /* wait for any signal to be caught */

/* some signal has been caught, SIGALR is now blocked */

unslept = alarm(0);
sigaction(SIGALRM, &oldact, NULL); /* reset previous action */

/* reset signal mask, which unblocks SIGALRM */
sigprocmask(SIG_SETMASK, &oldmasks, NULL);

return(unslept);
}
```

- Handles signals reliably
- Avoiding the race condition
- Do not handle any interactions with previously set alarms

# Job-Control Signals

SIGCHLD	: Child process has stopped or terminated
SIGCONT	: Continue process, if stopped
SIGSTOP	: Stop signal (can't be caught or ignored)
SIGTSTP	: Interactive stop signal
SIGTTIN	: Read from controlling terminal by a background process group
SIGTTOUT	: Write to controlling terminal by member of a background process group

- Most applications don't handle these signals - interactive shell do all the work required to handle these signals.
  - A program that manages the terminal needs to handle job-control signals (ex: vi editor – suspended when control-Z and redraws the screen when fg)
- When type Control-Z (suspend Character), SIGTSTP sent to all processes in the foreground process group.



# Job-Control Signals

```
/* How to handle SIGTSTP (handling a job-control signal) */
int main(void) {
    int n;char          buf[BUFSIZE];
    /* only catch SIGTSTP if we're running with a job-control shell */
    if (signal(SIGTSTP, SIG_IGN) == SIG_DFL)
        signal(SIGTSTP, sig_tstp);
    while ( (n = read(STDIN_FILENO, buf, BUFSIZE)) > 0)
        if (write(STDOUT_FILENO, buf, n) != n) err_sys("write error");
        if (n < 0) err_sys("read error");
    exit(0);
}

static void sig_tstp(int signo) {          /* signal handler for SIGTSTP */
    sigset_t  mask;
    /* ... move cursor to lower left corner, reset tty mode ... */

    /* unblock SIGTSTP, since it's blocked while we're handling it */
    sigemptyset(&mask);
    sigaddset(&mask, SIGTSTP);
    sigprocmask(SIG_UNBLOCK, &mask, NULL);

    signal(SIGTSTP, SIG_DFL);              /* reset disposition to default */

    kill(getpid(), SIGTSTP);               /* and send the signal to ourself */

        /* we won't return from the kill until we're continued */

    signal(SIGTSTP, sig_tstp);             /* reestablish signal handler */

    /* ... reset tty mode, redraw screen ... */
    return;
}
```